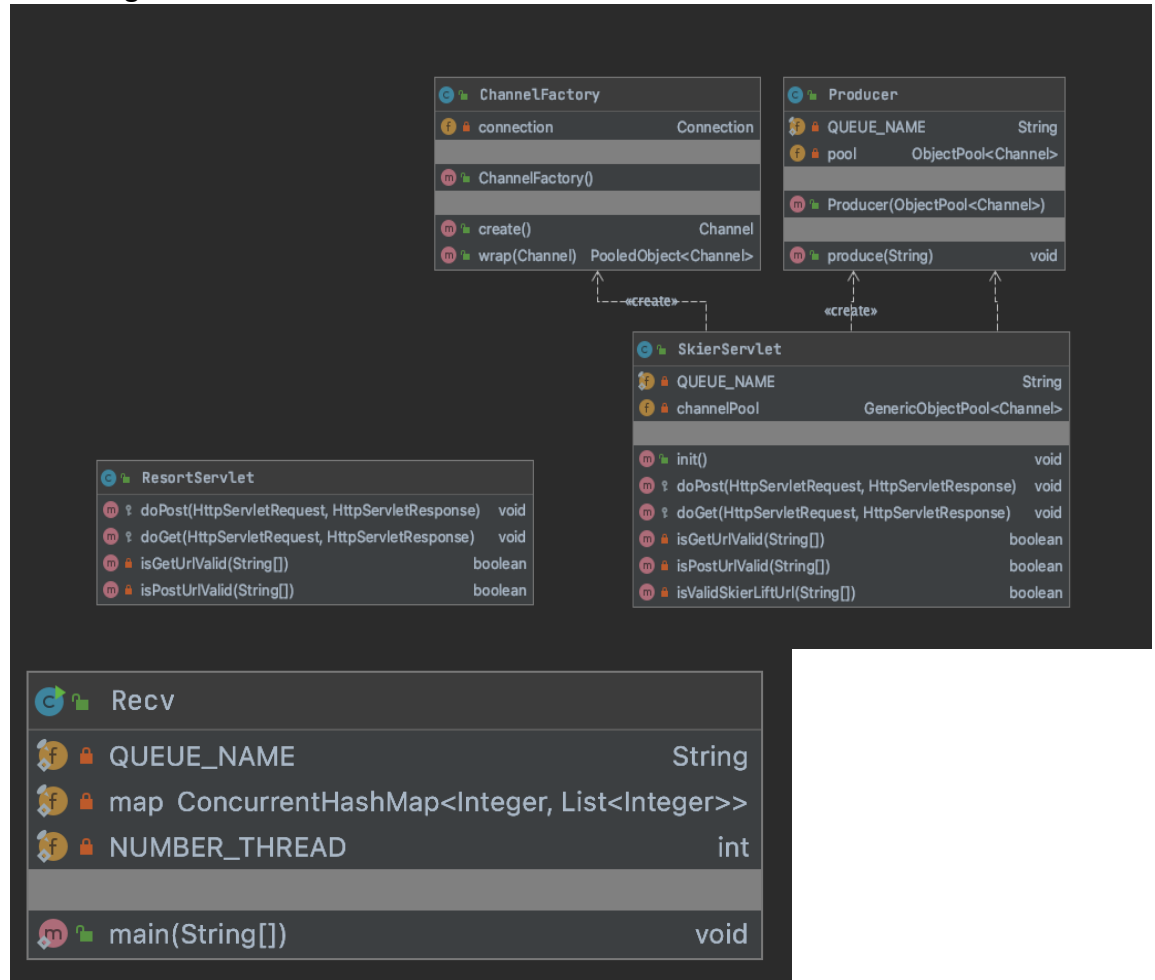


Assignment 2 – CS6650

Yuchen Tian

1. URL for Github Repo: <https://github.com/tyuchn/cs6650-hw2-ytian>
2. Server Design
 - a. Class Diagram



- b. Mechanisms

The server consists of a tomcat servlet (including a producer), a RabbitMQ message queue and a consumer. For the producer, since we mainly focus on the skier servlet. There are three main classes in the skier servlet: *SkierServlet*, *ChannelFactory* and *Producer*. The *SkierServlet* is refactored from the previous assignment, it has an instance of *GenericObjectPool* (in the *init()* method), which is initiated by a *ChannelFactory* class. This pool will serve as a channel pool that shares pre-created channels with the producers. The *ChannelFactory* class is a factory for channel creation, it is implemented on top of the apache pool2 package. *Producer* class is the producer for sending message to RabbitMQ, it is instantiated by a channel pool, and in the *produce()* method, it will borrow a channel object from the channel pool, then publish the message to the message queue, at last, it will return the borrowed channel to the channel pool for future

use. Recv is the consumer class, it was a 256-thread program, each thread creates a channel, get the message and writes to a concurrent hashmap. It is always implemented on the rabbitmq client package.

c. The life cycle of a message

When a post request was sent to the server, the skier servlet will first validate it, if the validation was passed, it will instantiate a producer object, and assign a channel from the channel pool to the producer, then the producer will publish the message to the message queue. On the consumer side, there is a 256-thread program, each thread will have a channel connecting with the message queue and consumer the message. At the same time, I use a concurrent map to store the record for each liftID.

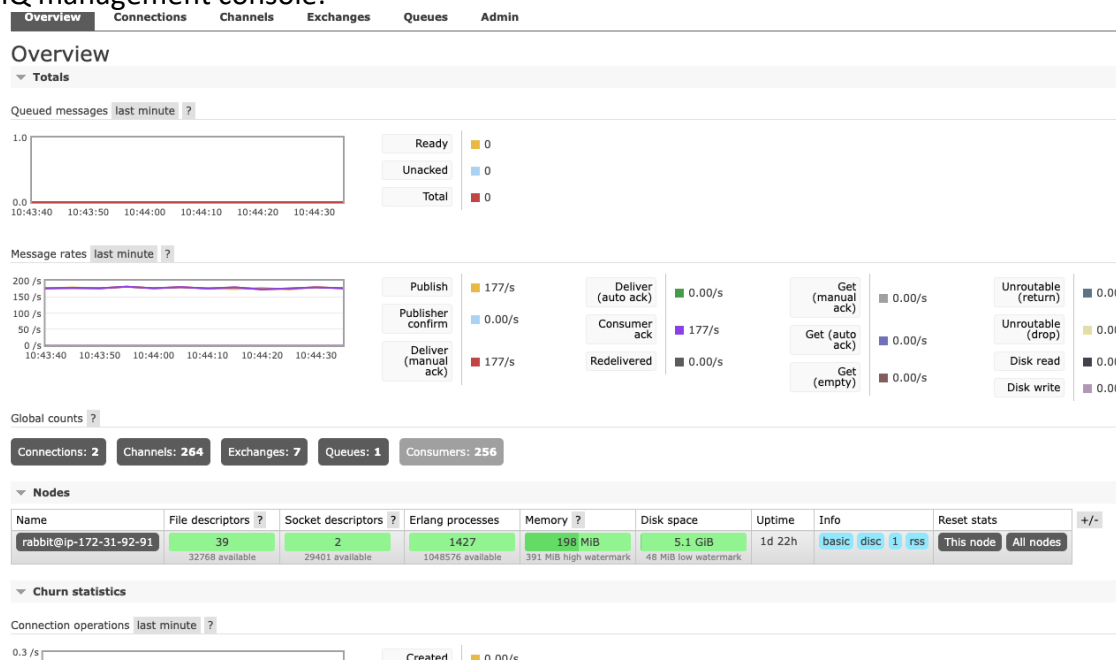
3. Test Results

a. 64 Threads

i. Commands lines:

`-numThreads 64 -numSkiers 100000 -address 54.82.5.13:8080`

ii. RMQ management console:



iii. Result:

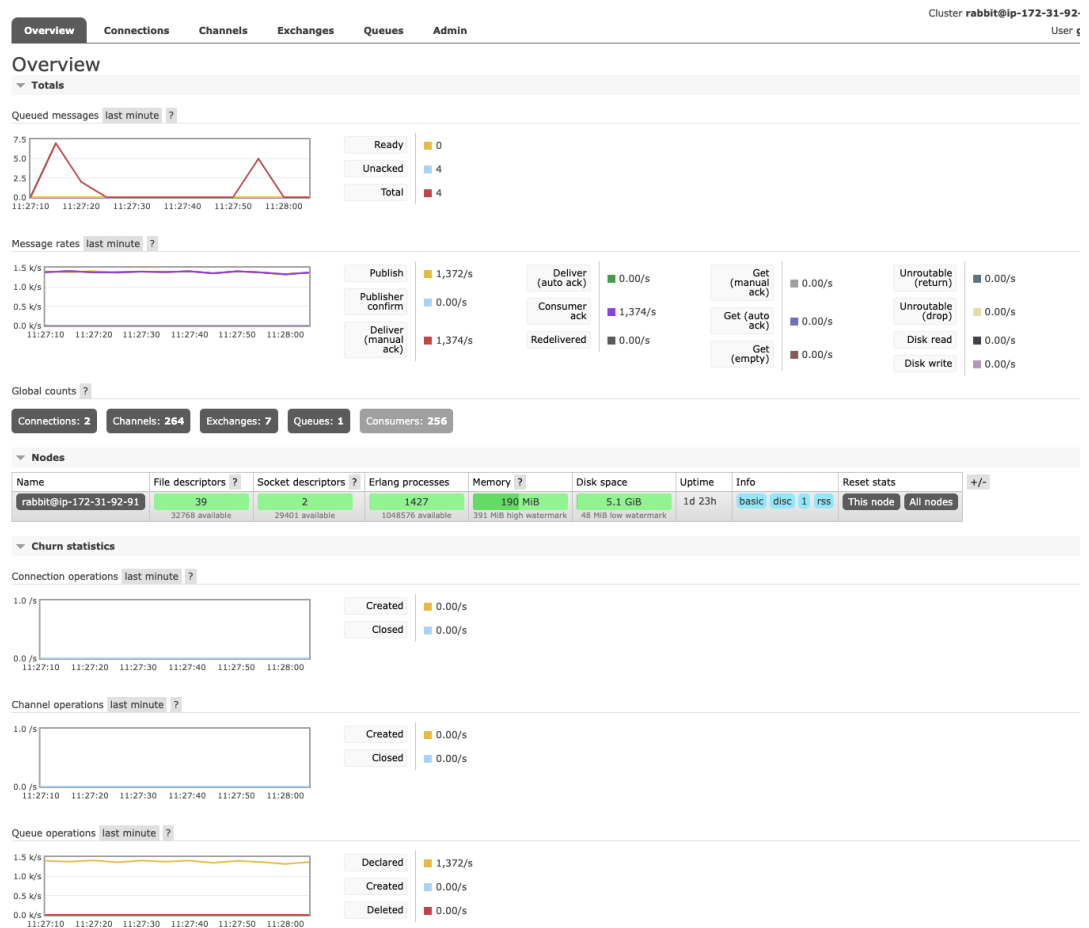
```
SkierClient x
/Library/Java/JavaVirtualMachines/temurin-11.jdk/Contents/Home/bin/java ...
Successful requests sent 799824
Unsuccessful requests 0
WallTime 1999692ms
Total throughput per ms 399.9736
|
Process finished with exit code 0
```

b. 128 Threads

i. Commands lines:

`-numThreads 128 -numSkiers 100000 -address 54.82.5.13:8080`

ii. RMQ management console:



iii. Result:

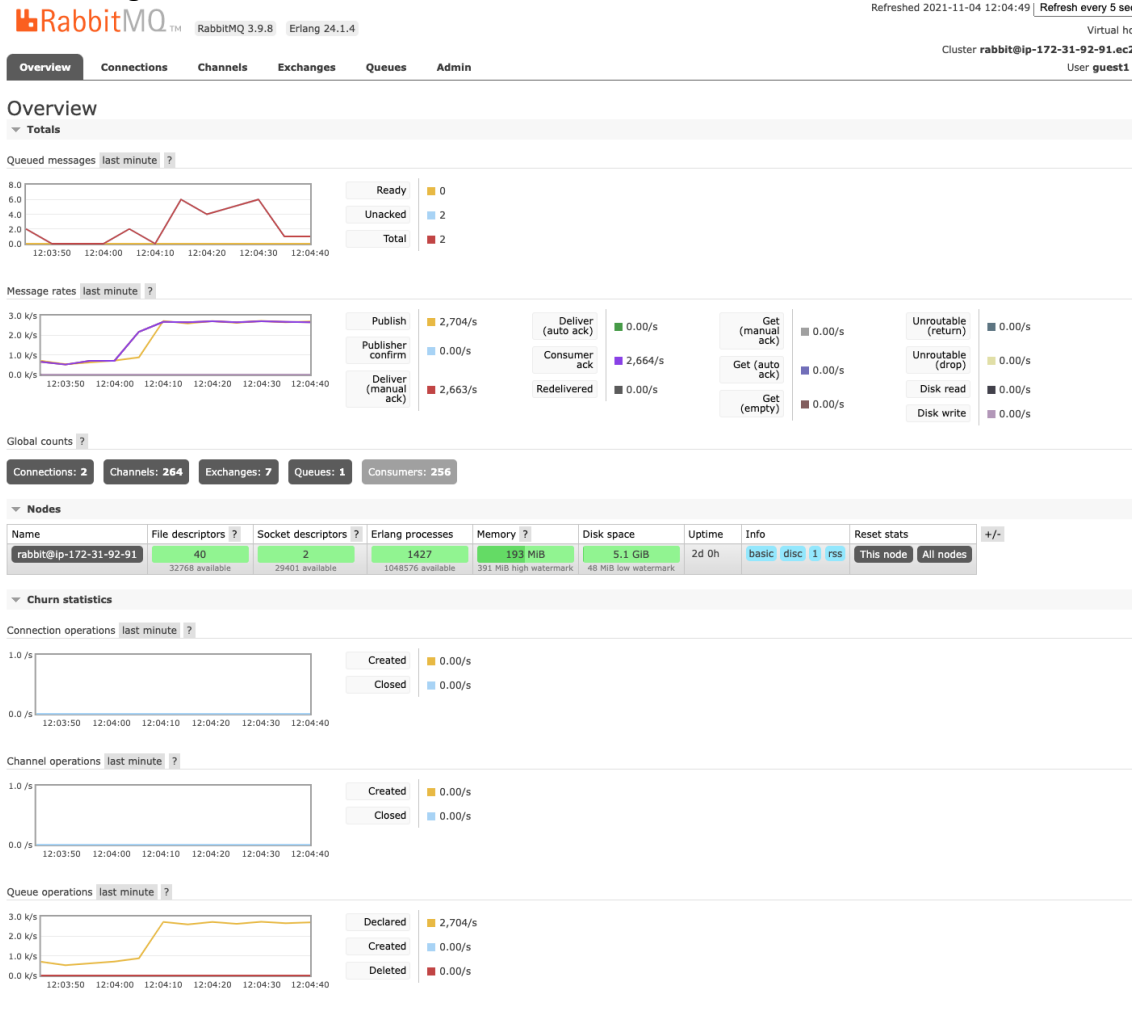
```
SkierClient x
/Library/Java/JavaVirtualMachines/tem
Successful requests sent 799840
Unsuccessful requests 0
WallTime 1007452ms
Total throughput per ms 793.9237
Process finished with exit code 0
```

c. 256 Threads

i. Commands lines:

-numThreads 256 -numSkiers 100000 -address 54.82.5.13:8080

ii. RMQ management console:



iii. Result:

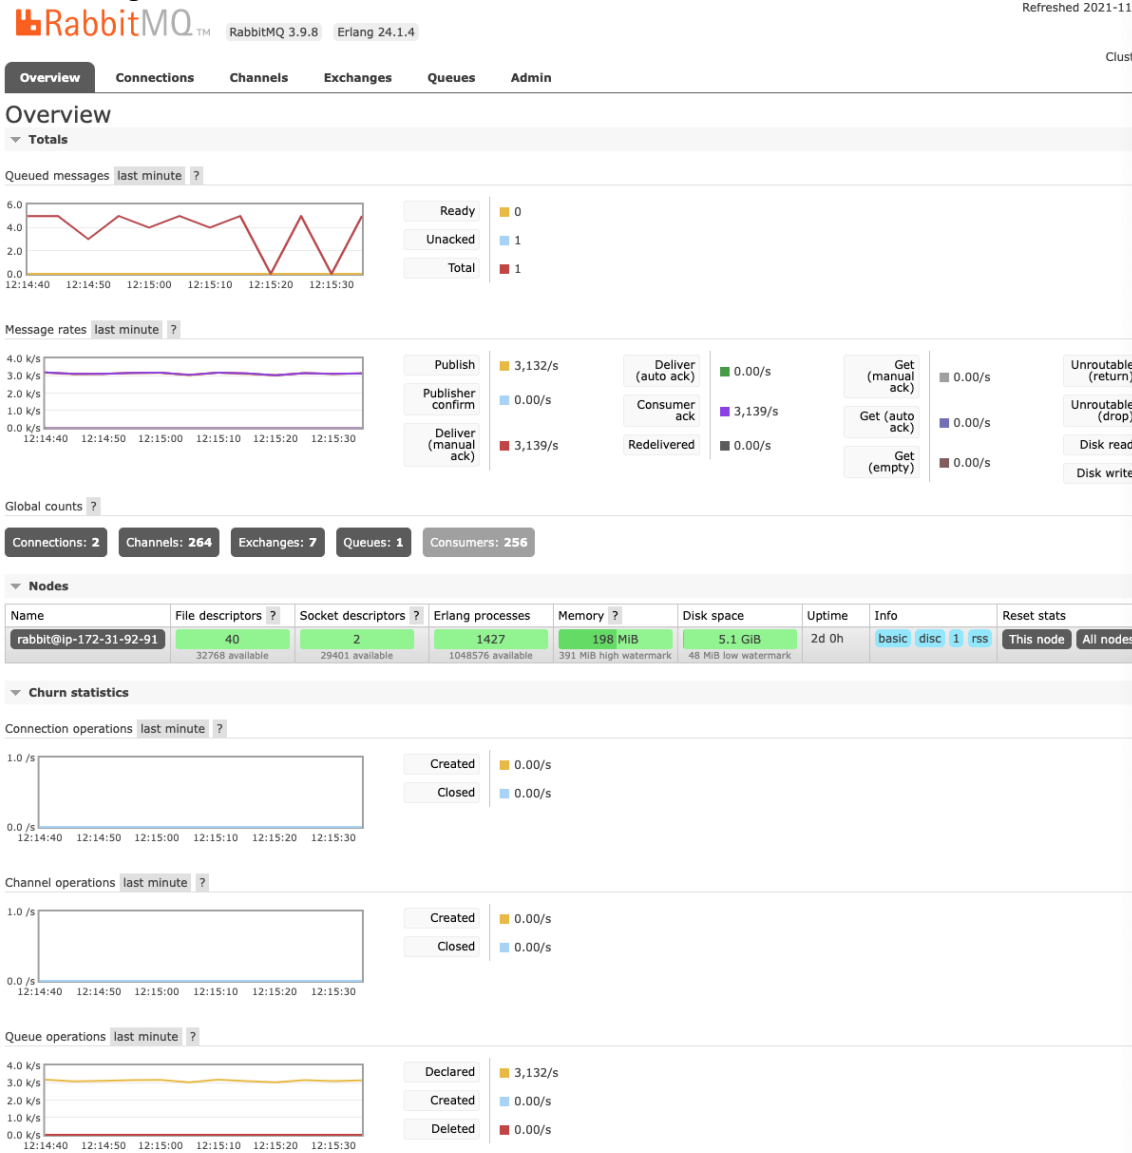
```
SkierClient x
/Library/Java/JavaVirtualMachines/tem
Successful requests sent 799040
Unsuccessful requests 0
WallTime 529197ms
Total throughput per ms 1509.9104
Process finished with exit code 0
```

d. 512 Threads

i. Commands lines:

-numThreads 512 -numSkiers 100000 -address 54.82.5.13:8080

ii. RMQ management console:



iii. Result:

```
SkierClient x
/Library/Java/JavaVirtualMachines/tem
Successful requests sent 799104
Unsuccessful requests 0
WallTime 341676ms|
Total throughput per ms 2338.777

Process finished with exit code 0
```

4. Result Analysis:

Based on the result, in most of the time, the message queue size is close to zero. Which means our 256-thread consumer is capable to consumer all the message sent by 64, 128, 256 and 512 threads producer, there's barely any message remains in the message queue. When a message was sent by the client, the producer will immediately send it to the queue, and there must be a consumer thread taking the message from its own channel at almost the same time. The testing result also indicates our design is efficient enough.